

Kothe Viswanath

+1-680-356-8964 | viswanathkothe@gmail.com | <https://linkedin.com/in/kothe-viswanath> | San Francisco, CA

EDUCATION

Syracuse University

Master of Science in Computer Engineering

Syracuse, NY

Aug 2023 – May 2025

EXPERIENCE

Software Engineer | Research Platform • Syracuse University • Remote, California Jun 2025 – Present

- Designed and operated a **multi-tenant GPU training platform** (SLURM) supporting 50+ researchers with distributed training, checkpointing, and fair-share scheduling across shared clusters
- Built **LLM inference stack** (vLLM + QLoRA) with production serving workflows enabling high-throughput experimentation and significantly faster model iteration across teams
- Orchestrated **500+ ML pipelines/day on AWS EKS** with sandbox isolation, quotas, and automated recovery
- Deployed **GitOps IaC and OpenTelemetry observability** to reducing job failures and diagnosis time by 60%

Software Engineer | AI/ML Platform • ISRO • Bangalore, India Jan 2022 – June 2023

- Owned **end-to-end ML infrastructure** across SLURM, Kubernetes, SAN/Infiniband storage, and **deployment systems**, supporting 50+ scientists running mission-critical satellite workloads
- Architected Kafka + Flink/Spark **telemetry pipelines** processing 1M+ events/day from **satellite edge systems**, enabling sub-200ms real-time analytics, monitoring, and downstream inference services
- Built Kubernetes-based **CI/CD deployment platform with canaries and rollbacks**, reducing model releases from days to under 30 minutes
- Deployed **real-time anomaly detection services (0.95 F1)** enabling proactive satellite failure alerts
- Developed **cluster-wide ML observability** for drift tracking, performance metrics, and GPU utilization

Software Engineer | Observability • Nuviso • Bangalore, India June 2020 – Dec 2021

- Architected **cloud-native streaming and monitoring platform** ingesting telemetry from 500+ devices into Kafka, enabling near real-time analytics and production visibility for SOC and operations teams
- Built **Go + OpenSearch + React dashboards** for incident response and system health monitoring
- Standardized **containerized CI/CD + GitOps workflows** across services and microservices

PERSONAL PROJECTS

Personal Agent • github.com/vishwa419/PersonalAgent Oct 2025 – Present

- Built AI-native browser assistant using GPT and MCP agent architecture for automated task tracking, deadline management, email workflows, and content summarization

Agentic Docker • github.com/vishwa419/Kitty_containers Aug 2025 – Sep 2025

- Built modular agentic-AI system for task orchestration integrating GPT, MCP, and RAG to automate container deployments with LangGraph, LangChain, and Qdrant

MLOps Pipeline • github.com/vishwa419/MLOps_Inference Jun 2025 – Aug 2025

- Built GitOps-driven ML pipeline with automated data validation, feature store, model training, deployment, and monitoring using BentoML, MLFlow, Feast, and Grafana

SKILLS

Languages: Go, Python, JavaScript

ML/AI Infrastructure: SLURM, vLLM, Kubeflow, PyTorch, TensorFlow

Cloud & Orchestration: Kubernetes (EKS), Terraform, ArgoCD, Docker

Data & Streaming: Kafka, Flink, Spark, OpenSearch, Redis

Observability: OpenTelemetry, Prometheus, Grafana, CloudWatch

Networking & Protocols: gRPC, REST, WebSocket, InfiniBand